

Importing/export, basic plotting

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Another very common way of "creating" a Pandas Dataframe is by importing a table from another format like CSV or Excel.

Simple import

An Excel table is provided in [composers.xlsx \(composers.xlsx\)](#) and can be read with the `read_excel` function. There are many more readers for other types of data (csv, json, html etc.) but we focus here on Excel.

```
In [2]: pd.read_excel('composers.xlsx')
```

Out[2]:

	composer	birth	death	city
0	Mahler	1860	1911	Kaliste
1	Beethoven	1770	1827	Bonn
2	Puccini	1858	1924	Lucques
3	Shostakovich	1906	1975	Saint-Petersburg

The reader automatically recognized the headers of the file. However it created a new index. If needed we can specify which column to use as header:

```
In [4]: pd.read_excel('composers.xlsx', index_col = 'composer')
```

Out[4]:

	birth	death	city
composer			
Mahler	1860	1911	Kaliste
Beethoven	1770	1827	Bonn
Puccini	1858	1924	Lucques
Shostakovich	1906	1975	Saint-Petersburg

If we open the file in Excel, we see that it is composed of more than one sheet. Clearly, when not specifying anything, the reader only reads the first sheet. However we can specify a sheet:

```
In [5]: specific_sheet = pd.read_excel('composers.xlsx', index_col = 'composer', sheet_
```

```
In [6]: specific_sheet
```

Out[6]:

	birth	death	city
composer			
Mahler	1860.0	1911	Kaliste
Beethoven	1770.0	1827	Bonn
Puccini	1858.0	1924	Lucques
Shostakovich	1906.0	1975	Saint-Petersburg
Sibelius	10.0	unknown	unknown
Haydn	NaN	NaN	Röhräu

For each reader, there is a long list of options to specify how the file should be read. We can see all these options using the help (see below). Imagine that our tables contains a title and unnecessary rows: we can use the `skiprows` argument. Imagine you have dates in your table: you can use the `date_parser` argument to specify how to format them etc.

```
In [7]: #use shift+tab within the parenthesis to see optional arguemnts  
# pd.read_excel()
```

Handling unknown values

As you can see above, some information is missing. Some missing values are marked as "unknown" while other are NaN. NaN is the standard symbol for unknown/missing values and is understood by Pandas while "unknown" is just seen as text. This is impractical as now we have e.g. columns with a mix of numbers and text which will make later computations difficult. What we would like to do is to replace all "irrelevant" values with the standard NaN symbol that says "no information".

Let's first do a regular import:

```
In [8]: import1 = pd.read_excel('composers.xlsx', index_col = 'composer', sheet_name='S')
import1
```

Out[8]:

	birth	death	city
composer			
Mahler	1860.0	1911	Kaliste
Beethoven	1770.0	1827	Bonn
Puccini	1858.0	1924	Lucques
Shostakovich	1906.0	1975	Saint-Petersburg
Sibelius	10.0	unknown	unknown
Haydn	NaN	NaN	Röhräu

If we look now at one column, we can see that columns have been imported in different ways. One column is an object, i.e. mixed types, the other contains floats:

```
In [9]: import1.birth
```

Out[9]:

composer	
Mahler	1860.0
Beethoven	1770.0
Puccini	1858.0
Shostakovich	1906.0
Sibelius	10.0
Haydn	NaN

Name: birth, dtype: float64

```
In [10]: import1.death
```

Out[10]:

composer	
Mahler	1911
Beethoven	1827
Puccini	1924
Shostakovich	1975
Sibelius	unknown
Haydn	NaN

Name: death, dtype: object

If we want to do calculations, for example getting summary information using `describe()` we have a problem: the `death` column is skipped because no calculation can be done with strings:

```
In [11]: import1.describe()
```

Out[11]:

	birth
count	5.000000
mean	1480.800000
std	823.674207
min	10.000000
25%	1770.000000
50%	1858.000000
75%	1860.000000
max	1906.000000

Now we specify that 'unknown' should be a NaN value:

```
In [12]: import2 = pd.read_excel('composers.xlsx', index_col = 'composer',
                                sheet_name='Sheet2', na_values=['unknown'])
import2
```

Out[12]:

	birth	death	city
composer			
Mahler	1860.0	1911.0	Kaliste
Beethoven	1770.0	1827.0	Bonn
Puccini	1858.0	1924.0	Lucques
Shostakovich	1906.0	1975.0	Saint-Petersburg
Sibelius	10.0	NaN	NaN
Haydn	NaN	NaN	Röhräu

And now computations are again possible, as Pandas knows how to deal with NaNs:

```
In [13]: import2.describe()
```

Out[13]:

	birth	death
count	5.000000	4.000000
mean	1480.800000	1909.250000
std	823.674207	61.396933
min	10.000000	1827.000000
25%	1770.000000	1890.000000
50%	1858.000000	1917.500000
75%	1860.000000	1936.750000
max	1906.000000	1975.000000

Handling bad or missing values is a very important part of data science. Taking care of the most common occurrences at import is a good solution.

Column types

We see above that the birth column has been "classified" as a float. However we know that this is not the case, it's just an integer. Here again, we can specify the column type already at import time using the dtype option and a dictionary:

```
In [12]: import2 = pd.read_excel('composers.xlsx', index_col = 'composer', sheet_name='S',
                                dtype={'composer':np.str, 'birth':np.int32, 'death':np.in
```

```
C:\Users\user\AppData\Local\Temp\ipykernel_7384\1656518521.py:2: DeprecationWarning: `np.str` is a deprecated alias for the builtin `str`. To silence this warning, use `str` by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.str_` here.
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
dtype={'composer':np.str, 'birth':np.int32, 'death':np.int32, 'city':np.str})
```

```
In [14]: import2.birth
```

Out[14]:

composer	
Mahler	1860.0
Beethoven	1770.0
Puccini	1858.0
Shostakovich	1906.0
Sibelius	10.0
Haydn	NaN

Name: birth, dtype: float64

Modifications after import

Of course we don't have to do all these adjustment at import time. We can also do a default import and check what has to be corrected afterward.

Create NaNs

If we missed some bad values at import we can just replace all those directly in the dataframe. We can achieve that by using the `replace()` method and specifying what should be replaced:

```
In [15]: import1
```

```
Out[15]:
```

	birth	death	city
composer			
Mahler	1860.0	1911	Kaliste
Beethoven	1770.0	1827	Bonn
Puccini	1858.0	1924	Lucques
Shostakovich	1906.0	1975	Saint-Petersburg
Sibelius	10.0	unknown	unknown
Haydn	NaN	NaN	Röhräu

```
In [16]: import_nans = import1.replace('unknown', np.nan)
import_nans.birth
```

```
Out[16]: composer
Mahler      1860.0
Beethoven   1770.0
Puccini     1858.0
Shostakovich 1906.0
Sibelius    10.0
Haydn       NaN
Name: birth, dtype: float64
```

Note that when we fix "bad" values, e.g. here the "unknown" text value with NaNs, Pandas automatically adjust the type of the column, allowing us for example to later do mathematical operations.

```
In [17]: import1.death.dtype
```

```
Out[17]: dtype('O')
```

```
In [18]: import_nans.death.dtype
```

```
Out[18]: dtype('float64')
```

Changing the type

We can also change the type of a column on an existing Dataframe with the same command as in Numpy:

```
In [19]: import2.birth
```

```
Out[19]: composer
Mahler      1860.0
Beethoven   1770.0
Puccini      1858.0
Shostakovich 1906.0
Sibelius     10.0
Haydn        NaN
Name: birth, dtype: float64
```

```
In [20]: import2.birth.astype('float')
```

```
Out[20]: composer
Mahler      1860.0
Beethoven   1770.0
Puccini      1858.0
Shostakovich 1906.0
Sibelius     10.0
Haydn        NaN
Name: birth, dtype: float64
```

If we look again at import2:

```
In [21]: import2.birth
```

```
Out[21]: composer
Mahler      1860.0
Beethoven   1770.0
Puccini      1858.0
Shostakovich 1906.0
Sibelius     10.0
Haydn        NaN
Name: birth, dtype: float64
```

we see that we didn't actually change the type. Changes on a Dataframe are only effective if we reassign the column:

```
In [22]: import2.birth = import2.birth.astype('float')
```

```
In [23]: import2.birth
```

```
Out[23]: composer
Mahler      1860.0
Beethoven   1770.0
Puccini     1858.0
Shostakovich 1906.0
Sibelius    10.0
Haydn       NaN
Name: birth, dtype: float64
```

Export

You can easily export a Dataframe that you worked on. Most commonly you will export it in a common format like CSV:

```
In [24]: import2.to_csv('mydataframe.csv')
```

If you have a complex dataframe that e.g. contains lists, you can save it as a *pickle* object, a specific Python format that allows one to save complex data:

```
In [25]: import2.to_pickle('my_dataframe.pkl')
```

You can reload this type of data via the pickle loading function of Pandas:

```
In [26]: import3 = pd.read_pickle('my_dataframe.pkl')
```

```
In [27]: import3
```

```
Out[27]:
```

	birth	death	city
composer			
Mahler	1860.0	1911.0	Kaliste
Beethoven	1770.0	1827.0	Bonn
Puccini	1858.0	1924.0	Lucques
Shostakovich	1906.0	1975.0	Saint-Petersburg
Sibelius	10.0	NaN	NaN
Haydn	NaN	NaN	Röhräu


Plotting

We will learn more about plotting later, but let's see here some possibilities offered by Pandas. Pandas builds on top of Matplotlib but exploits the knowledge included in Dataframes to improve the default output. Let's see with a simple dataset.


```
In [28]: composers = pd.read_excel('composers.xlsx', sheet_name='Sheet5')
```


We can pass Series to Matplotlib which manages to understand them. Here's a default scatter plot:

```
In [29]: plt.plot(composers.birth, composers.death, 'o')  
plt.show()
```



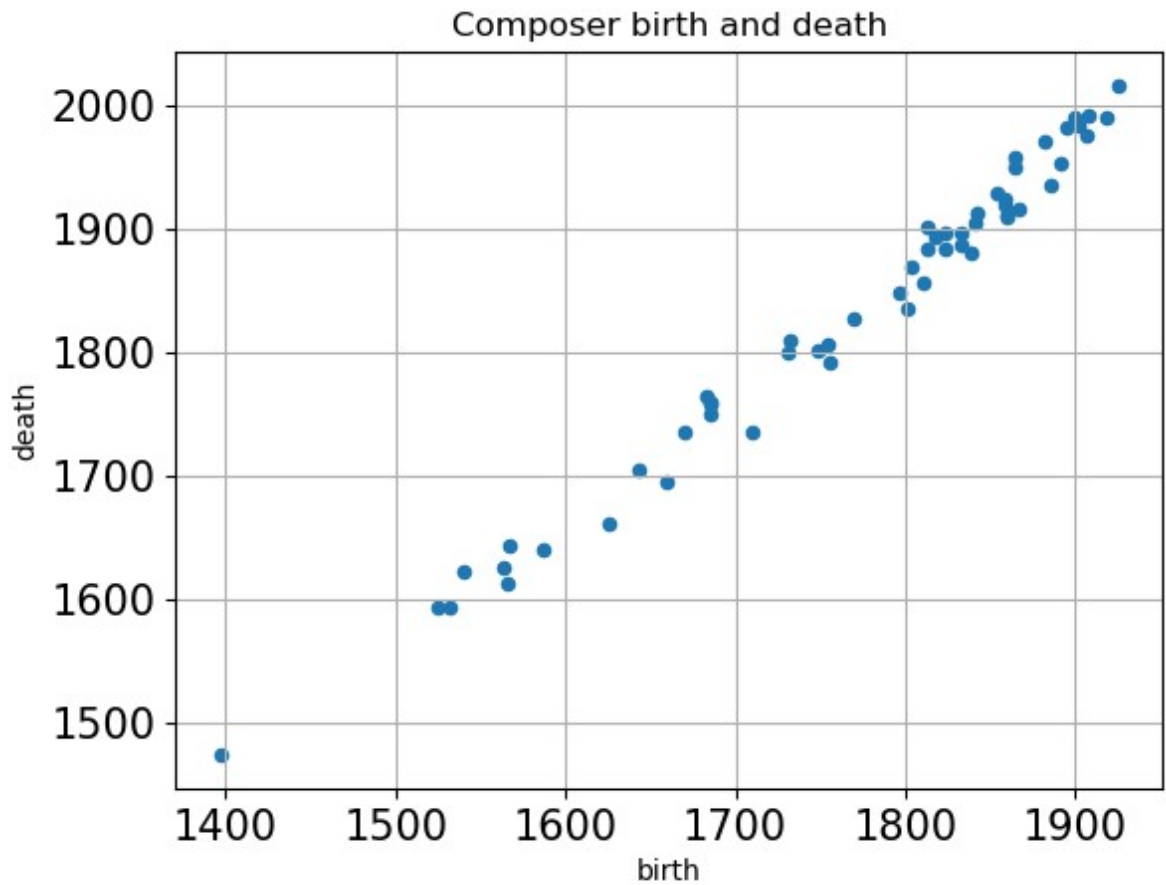
Now we look at the default Pandas output. Different types of plots are accessible when using the `data_frame.plot` function via the `kind` option. The variables to plot are column names passed as keywords instead of whole series like in Matplotlib:

```
In [30]: composers.plot(x = 'birth', y = 'death', kind = 'scatter')
plt.show()
```




We see that the plot automatically gets axis labels. Another gain is that some obvious options like setting a title are directly accesible when creating the plot:

```
In [31]: composers.plot(x = 'birth', y = 'death', kind = 'scatter',  
                        title = 'Composer birth and death', grid = True, fontsize = 15)  
plt.show()
```



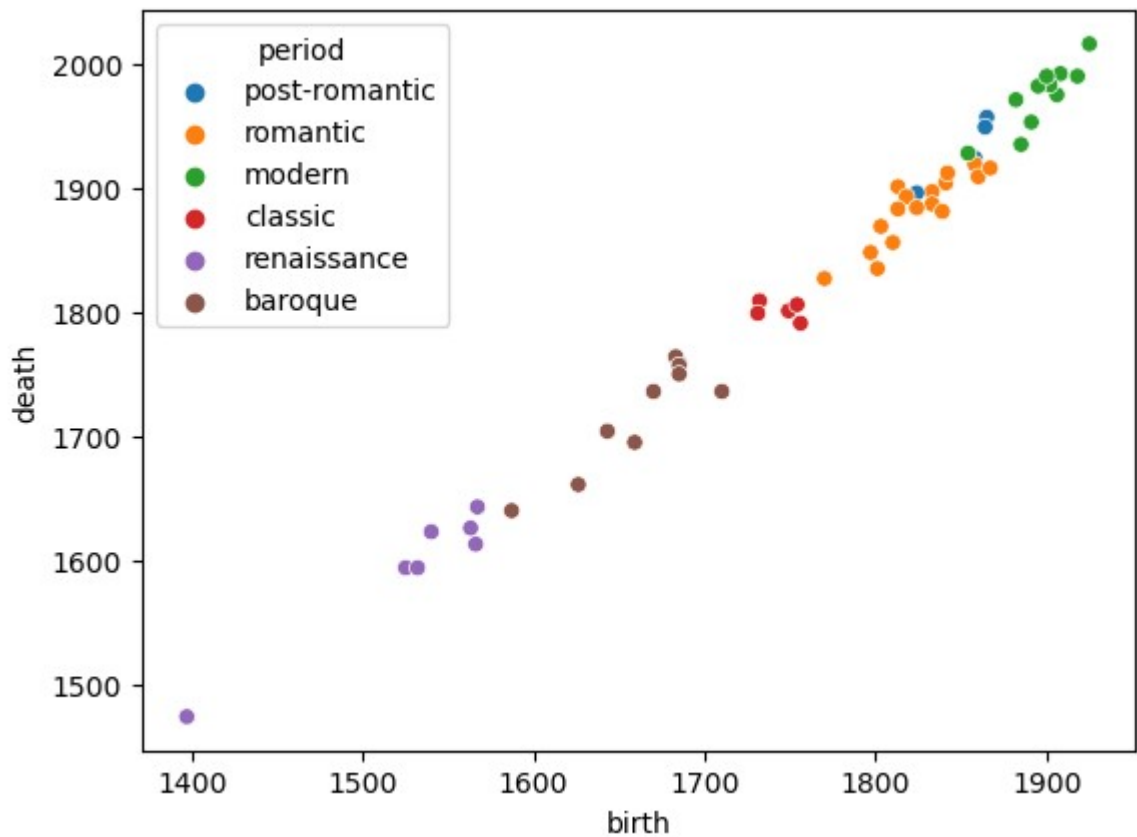
One can add even more information on the plot by using more arguments used in a similar way as a grammar of graphics. For example we can color the scatter plot by periods:

```
In [32]: composers.plot(x = 'birth', y = 'death', kind = 'scatter',  
                        c = composers.period.astype('category').cat.codes, colormap = '  
plt.show()
```



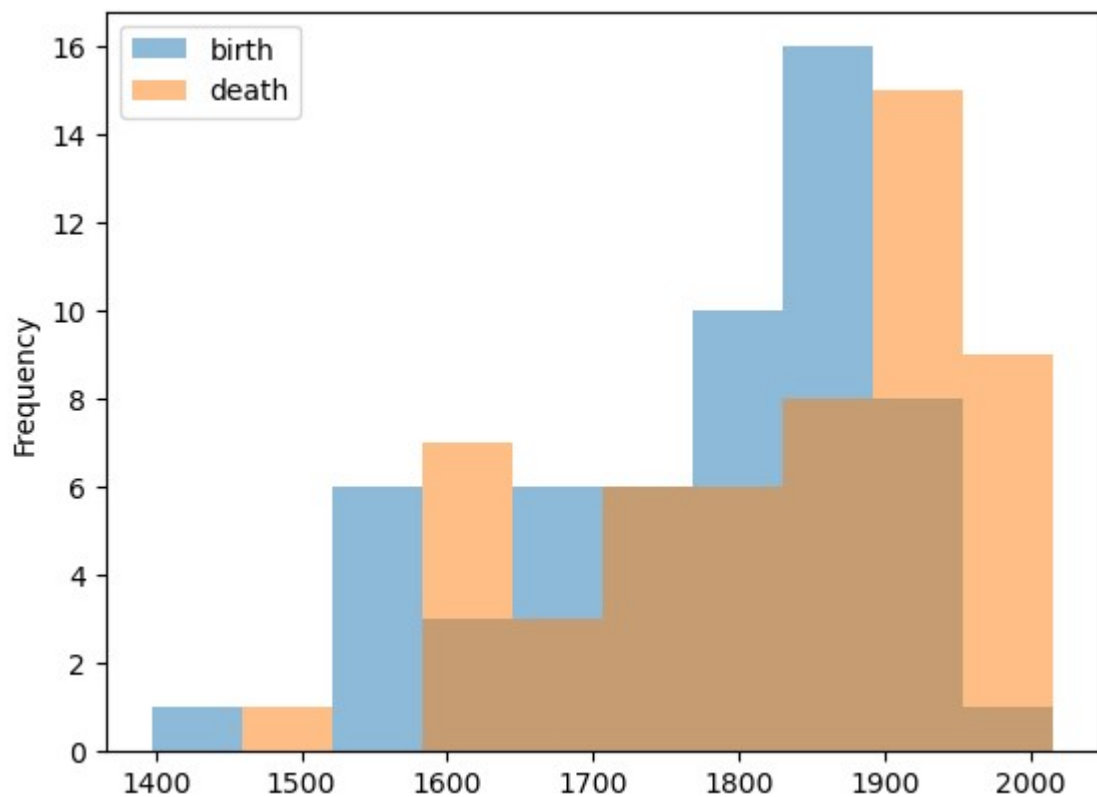
Here you see already a limitation of the plotting library. To color dots by the period category, we had to turn the latter into a series of numbers. We could then rename those to improve the plot, but it's better to use more specialized packages such as Seaborn which allow to realize this kind of plot easily:

```
In [33]: sns.scatterplot(data = composers, x = 'birth', y = 'death', hue = 'period')  
plt.show()
```



Some additional plotting options are available in the `plot()` module. For example histograms:

```
In [34]: composers.plot.hist(alpha = 0.5)
plt.show()
```



Here you see again the gain from using Pandas: without specifying anything, Pandas made a histogram of the two columns containing numbers, labelled the axis and even added a legend to the plot.

All these features are very nice and very helpful when exploring a dataset. When analyzing data in depth and creating complex plots, Pandas's plotting might however be limiting and other options such as Seaborn or Plotnine can be used.

Finally, all plots can be "styled" down to the smallest detail, either by using Matplotlib options or by directly applying a style e.g.:

```
In [35]: plt.style.use('ggplot')
```

```
In [36]: composers.plot.hist(alpha = 0.5)  
plt.show()
```

